

```
In [1]: import pandas as pd
import statistics as st
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv('titanic (1).csv')
df.head(2)
df.tail(2)
```

Out[2]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.00	C148	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.75	NaN	Q

```
In [3]: df.shape
```

Out[3]: (891, 12)

Exploratory Data Analysis

```
In [4]: df.dtypes
```

```
Out[4]: PassengerId    int64
Survived              int64
Pclass               int64
Name                 object
Sex                  object
Age                 float64
SibSp                int64
Parch                int64
Ticket              object
Fare                 float64
Cabin                object
Embarked            object
dtype: object
```

In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age         714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [6]: `df.describe(include='all')`

Out[6]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
count	891.000000	891.000000	891.000000	891	891	714.000000	891.000000	891.000000	891	891.000000	204	889
unique	NaN	NaN	NaN	891	2	NaN	NaN	NaN	681	NaN	147	3
top	NaN	NaN	NaN	Braund, Mr. Owen Harris	male	NaN	NaN	NaN	347082	NaN	B96 B98	S
freq	NaN	NaN	NaN	1	577	NaN	NaN	NaN	7	NaN	4	644
mean	446.000000	0.383838	2.308642	NaN	NaN	29.699118	0.523008	0.381594	NaN	32.204208	NaN	NaN
std	257.353842	0.486592	0.836071	NaN	NaN	14.526497	1.102743	0.806057	NaN	49.693429	NaN	NaN
min	1.000000	0.000000	1.000000	NaN	NaN	0.420000	0.000000	0.000000	NaN	0.000000	NaN	NaN
25%	223.500000	0.000000	2.000000	NaN	NaN	20.125000	0.000000	0.000000	NaN	7.910400	NaN	NaN
50%	446.000000	0.000000	3.000000	NaN	NaN	28.000000	0.000000	0.000000	NaN	14.454200	NaN	NaN
75%	668.500000	1.000000	3.000000	NaN	NaN	38.000000	1.000000	0.000000	NaN	31.000000	NaN	NaN
max	891.000000	1.000000	3.000000	NaN	NaN	80.000000	8.000000	6.000000	NaN	512.329200	NaN	NaN

In [7]: `df.corr()`

Out[7]:


	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
PassengerId	1.000000	-0.005007	-0.035144	0.036847	-0.057527	-0.001652	0.012658
Survived	-0.005007	1.000000	-0.338481	-0.077221	-0.035322	0.081629	0.257307
Pclass	-0.035144	-0.338481	1.000000	-0.369226	0.083081	0.018443	-0.549500
Age	0.036847	-0.077221	-0.369226	1.000000	-0.308247	-0.189119	0.096067
SibSp	-0.057527	-0.035322	0.083081	-0.308247	1.000000	0.414838	0.159651
Parch	-0.001652	0.081629	0.018443	-0.189119	0.414838	1.000000	0.216225
Fare	0.012658	0.257307	-0.549500	0.096067	0.159651	0.216225	1.000000

```
In [8]: df.isnull().sum()
```

```
Out[8]: PassengerId      0  
Survived      0  
Pclass      0  
Name      0  
Sex      0  
Age      177  
SibSp      0  
Parch      0  
Ticket      0  
Fare      0  
Cabin      687  
Embarked      2  
dtype: int64
```

```
In [9]: df['Age'].hist()
```


```
Out[9]: <AxesSubplot:>
```



```
In [10]: sns.pairplot(df)
```


```
Out[10]: <seaborn.axisgrid.PairGrid at 0x1e7e0506f50>
```






```
In [11]: sns.heatmap(df.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

```
Out[11]: <AxesSubplot:>
```




```
In [12]: sns.heatmap(df.corr())
```

```
Out[12]: <AxesSubplot:>
```




```
In [13]: sns.countplot(x='Survived',data=df)
```

```
Out[13]: <AxesSubplot:xlabel='Survived', ylabel='count'>
```




```
In [14]: sns.boxplot(x='Pclass', y='Age', data =df)
```

```
Out[14]: <AxesSubplot:xlabel='Pclass', ylabel='Age'>
```



data wrangling:- cleaning data

```
In [15]: # plt.scatter(df['Age'],df['Pclass'])
# plt.bar(df['Age'],df['Pclass'])
df.head()
```

Out[15]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```
In [16]: df.dropna(inplace=True)
df.isnull().sum()
```

```
Out[16]: PassengerId    0
Survived              0
Pclass               0
Name                 0
Sex                  0
Age                  0
SibSp                0
Parch                0
Ticket               0
Fare                 0
Cabin                0
Embarked             0
dtype: int64
```

In [17]: `df.dtypes`

Out[17]:

```

PassengerId    int64
Survived       int64
Pclass         int64
Name           object
Sex            object
Age           float64
SibSp          int64
Parch          int64
Ticket         object
Fare           float64
Cabin          object
Embarked       object
dtype: object

```

In [18]: `df.corr()`

Out[18]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
PassengerId	1.000000	0.148495	-0.089136	0.030933	-0.083488	-0.051454	0.029740
Survived	0.148495	1.000000	-0.034542	-0.254085	0.106346	0.023582	0.134241
Pclass	-0.089136	-0.034542	1.000000	-0.306514	-0.103592	0.047496	-0.315235
Age	0.030933	-0.254085	-0.306514	1.000000	-0.156162	-0.271271	-0.092424
SibSp	-0.083488	0.106346	-0.103592	-0.156162	1.000000	0.255346	0.286433
Parch	-0.051454	0.023582	0.047496	-0.271271	0.255346	1.000000	0.389740
Fare	0.029740	0.134241	-0.315235	-0.092424	0.286433	0.389740	1.000000

In [19]: `df.drop(['PassengerId', 'Ticket', 'Cabin', 'Name'], axis=1, inplace=True)`

In [20]: df

Out[20]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
1	1	1	female	38.0	1	0	71.2833	C
3	1	1	female	35.0	1	0	53.1000	S
6	0	1	male	54.0	0	0	51.8625	S
10	1	3	female	4.0	1	1	16.7000	S
11	1	1	female	58.0	0	0	26.5500	S
...
871	1	1	female	47.0	1	1	52.5542	S
872	0	1	male	33.0	0	0	5.0000	S
879	1	1	female	56.0	0	1	83.1583	C
887	1	1	female	19.0	0	0	30.0000	S
889	1	1	male	26.0	0	0	30.0000	C

183 rows × 8 columns

In [21]: df

Out[21]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
1	1	1	female	38.0	1	0	71.2833	C
3	1	1	female	35.0	1	0	53.1000	S
6	0	1	male	54.0	0	0	51.8625	S
10	1	3	female	4.0	1	1	16.7000	S
11	1	1	female	58.0	0	0	26.5500	S
...
871	1	1	female	47.0	1	1	52.5542	S
872	0	1	male	33.0	0	0	5.0000	S
879	1	1	female	56.0	0	1	83.1583	C
887	1	1	female	19.0	0	0	30.0000	S
889	1	1	male	26.0	0	0	30.0000	C

183 rows × 8 columns

```
In [22]: import numpy as np
outlier = []
def detect_z(data):
    thres = 3
    mean = np.mean(data)
    std = np.std(data)

    for i in data:
        z = (i-mean)/std
        if (np.abs(z) > thres):
            outlier.append(i)
    print(outlier)

detect_z(df['Fare'])
```


```
[512.3292, 512.3292]
```

```
In [23]: sns.boxplot(df['Fare'])
```

C:\python314\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[23]: <AxesSubplot:xlabel='Fare'>
```




```
In [24]: for i in df.index:
         if df.loc[i, 'Fare'] > 150:
             df.drop(i, inplace=True)
```

```
In [25]: sns.boxplot(df['Fare'])
         print(df.shape)
```

(160, 8)

C:\python314\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



change datatype

```
In [26]: # df['Embarked'] = df['Embarked'].astype(int)
```

replace str with null

```
In [27]: # df['Embarked'] = df['Embarked'].replace('C', np.nan)
# df.dropna(inplace=True)
```

algorithm apply

```
In [28]: from sklearn.preprocessing import LabelEncoder
df1 = df.copy()
e1 = LabelEncoder()
e2 = LabelEncoder()

df1.Sex = e1.fit_transform(df1.Sex)
df1.Embarked = e2.fit_transform(df1.Embarked)

df1
```

Out[28]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
1	1	1	0	38.0	1	0	71.2833	0
3	1	1	0	35.0	1	0	53.1000	2
6	0	1	1	54.0	0	0	51.8625	2
10	1	3	0	4.0	1	1	16.7000	2
11	1	1	0	58.0	0	0	26.5500	2
...
871	1	1	0	47.0	1	1	52.5542	2
872	0	1	1	33.0	0	0	5.0000	2
879	1	1	0	56.0	0	1	83.1583	0
887	1	1	0	19.0	0	0	30.0000	2
889	1	1	1	26.0	0	0	30.0000	0

160 rows × 8 columns


```
In [29]: df = df1  
df.dtypes
```

```
Out[29]: Survived      int64  
Pclass      int64  
Sex         int32  
Age         float64  
SibSp       int64  
Parch       int64  
Fare        float64  
Embarked    int32  
dtype: object
```

```
In [30]: # X = df[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']]  
X = df.drop('Survived', axis=1)  
y = df['Survived']
```

```
In [31]: from sklearn.model_selection import train_test_split
```

```
In [32]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=.30, random_state=5)
```

```
In [33]: # from sklearn.linear_model import LinearRegression  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.linear_model import LogisticRegression
```

```
In [34]: T = LogisticRegression()
```

```
In [35]: T.fit(X_train,y_train)           #train dataset
```

C:\python314\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[35]: LogisticRegression()
```

```
In [36]: pred = T.predict(X_test)         #predict by      x test
```

```
In [37]: pred
```

```
Out[37]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1,
                1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1,
                1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1,
                1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0,
                0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1,
                1, 1], dtype=int64)
```

```
In [38]: from sklearn.metrics import classification_report,accuracy_score,confusion_matrix
```

```
In [39]: print(classification_report(y_test,pred))           #y test use
```

	precision	recall	f1-score	support
0	0.65	0.50	0.56	40
1	0.75	0.85	0.80	72
accuracy			0.72	112
macro avg	0.70	0.67	0.68	112
weighted avg	0.71	0.72	0.71	112

```
In [40]: print(accuracy_score(y_test,pred))
```

```
0.7232142857142857
```

```
In [41]: confusion_matrix(y_test,pred)
```

```
Out[41]: array([[20, 20],  
               [11, 61]], dtype=int64)
```

```
In [ ]:
```